

Bullet Farm: A Software-Change Transaction Processor for Heterogeneous Agents

Bullet Farm Maintainers

Abstract—Modern coding systems are commonly organized around an agent loop, a sandbox, a worktree, a reviewer, or a merge queue. Those mechanisms are useful but do not by themselves implement a software-change transaction. The missing contract binds who may mutate an exact repository subject, which process incarnation owns that authority, what evidence was obtained for which bytes, and how an ambiguous remote effect is reconciled without a blind retry. Bullet Farm is designed as a transaction processor around heterogeneous agents. Its binding chain is: incarnation-fenced authority, one repository writer, private Attempt clones, exact-OID Candidates, independently bound Evidence, durable effect intent, read-back and adoption after ambiguity, and truthful UNKNOWN projection across five trust planes. This paper evaluates the architecture against a dated public corpus: Gas Town v1.2.1, Gas City v1.4.1, DeepSeek Harness dsh-v0.1.1-rc.2, Omnigent v0.10.0, selected issue reports, and a broader tool landscape. It separates operational features from authority/integrity properties and uses four maturity labels: **Designed**, **Component-proved**, **Transaction-proved**, and **Live/release-proved**. At evidence snapshot 20260825-stage1-r3, Bullet Farm has component receipts but no connected TRANSACTION_PROOF, live proof, or release proof; the release inventory reports 26 blocked of 26 evaluated gates. Accordingly, this Stage-1 paper makes no performance or universal-superiority claim. A Stage-2 revision and matched benchmark are deferred until the connected transaction proof exists. The paper and its four-page executive brief consume one versioned JSON evidence manifest and are subject to deterministic publication preflight.

Index Terms—coding agents, transaction processing, Git isolation, evidence binding, authority fencing, effect reconciliation, merge safety.

I. RESEARCH CONTRACT

Agent systems have become capable enough that their limiting failure is often not model reasoning but the authority surrounding it. Two sessions can edit one checkout; a stale reviewer can approve bytes that no longer exist; a timed-out push can be retried after the server already applied it; and a UI can turn missing state into apparent success. Classical transaction-processing work treats identity, concurrency, logging, recovery, and idempotence as one problem [1]–[3]. Coding-agent stacks frequently distribute those obligations among an editor, Git, CI, a provider process, and a forge without an end-to-end subject binding.

Bullet Farm’s thesis is deliberately narrower than “better agents.” It is a software-change transaction processor *around* heterogeneous agents. A model may reason, search, and propose, but it cannot mint its own write, verification, or integration authority [4]. The research questions are:

- RQ1: Which identities and exact subjects must be bound so that a concurrent agent run can be recovered without inferring completion?
- RQ2: Which bindings are implemented and component-proved in the current family, and which remain design or connected-product gaps?
- RQ3: Do the reviewed pinned public contracts of adjacent systems document the same end-to-end bindings, after separating operational capability from authority and integrity?
- RQ4: What observations would falsify the design, and what matched experiment would be required before a comparative performance claim?

A. Inclusion and Source Discipline

Deep studies were included when a system exposed a public architecture plus a release or commit suitable for reproducible inspection. The fixed subjects are Gas Town v1.2.1, Gas City v1.4.1, DeepSeek Harness dsh-v0.1.1-rc.2, and Omnigent v0.10.0 [5]–[8]. Issue reports were included only when they described a concrete authority, Git, evidence, or recovery failure and could be mapped to a named invariant. The wider landscape appendix uses release/commit pins for open-source systems and official-document observation dates for proprietary ones. These are public-contract observations, not reverse engineering. Codex statements rely only on official OpenAI documentation [9], [10].

Every comparison cell uses exactly one of five terms. Documented means a pinned primary source describes the property. Partial/configuration-dependent means the property depends on configuration or covers only part of the row. Not documented means the reviewed sources do not describe it; it does not assert that the implementation lacks it. Unknown means the public material was insufficient to classify. N/A means the property is outside the product’s stated role. Absence of public documentation is never converted into a categorical “No.”

B. Evidence Vocabulary and Claim Limit

Table I is the paper’s claim lattice. The terms intentionally avoid the ambiguous word “complete.” They are generated into both documents from `evidence.json`, whose schema is `bullet.paper-evidence.v1` [11].

The snapshot records BLOCKED doctor status, PASS for the fast lane, PASS for the demo, and PASS for 2 bounded formal models. The release result is written as “26 blocked of 26 evaluated gates,” never as a bare fraction that could be mistaken for success. UNKNOWN, timeout, skipped, unsupported,

TABLE I
MATURITY LABELS USED THROUGHOUT THIS PAPER.

Label	Admitted meaning
Designed	Reviewed contract or invariant; no executable receipt required.
Component-proved	Exact committed component and mapped local receipt; no cross-process transaction inference.
Transaction-proved	One exact connected offline five-plane transaction with independent subject-bound receipts.
Live/release-proved	Admitted live provider/forged transaction, or signed release/install/recovery proof from the same subjects.

zero tests, a model assertion, HTTP 2xx, or process exit zero is not verified engineering truth [12]. No result from an active dirty tree is admitted into the committed evidence manifest.

II. THREAT AND FAILURE MODEL

The V1 threat model is local-first and single-user. It treats model output, repository content, concurrent agent incarnations, retrying supervisors, provider subprocesses, and network responses as untrusted. It trusts the host’s admitted binaries and kernel primitives within explicitly tested boundaries; remote multi-tenant runners and Byzantine forge operators are out of scope. The same local UID remains a material limitation because it can race path-based tools that have not been sealed or made descriptor-relative.

The recurring failures are not peculiar to agents. They are ordinary distributed-systems failures expressed through Git:

- 1) *Authority aliasing*. Two living processes believe they own the same Attempt or writable checkout. A session name is not an incarnation fence.
- 2) *Repository aliasing*. Linked worktrees have per-worktree HEAD and index, but share the object database, common metadata, and most refs [13]. They improve workspace ergonomics; they do not create separate repository principals.
- 3) *Subject drift*. Tests or reviews name a branch, workspace, or PR rather than the exact base, head, tree, patch, environment, and gate set. Rebase or merge can silently invalidate the result.
- 4) *Ambiguous effects*. A timeout occurs after a server may have accepted a push or workflow creation. Retrying before read-back can duplicate the logical mutation.
- 5) *Inferred liveness*. tmux, a PID, an idle terminal, or a model’s “done” statement becomes a completion oracle. Silence is then confused with death or success.
- 6) *Destructive recovery*. Cleanup removes a workspace before durable preservation and read-back establish whether the last mutation happened.
- 7) *Projection promotion*. A dashboard converts absent, stale, or contradictory ledger state into a green label. The UI becomes accidental authority.

The public issue corpus supplies concrete instances: town-root checkout deleted operational files [14]; parallel rig-root writers clobbered edits [15]; a timeout after a server commit admitted duplicate workflows [16]; and a non-atomic claim let two sessions execute the same task [17]. Inferred liveness has its own instances: a polecat session that died mid-task left

its work hooked with no notification [18]; a patrol cleanup step whose liveness check classified every worker as dead would have force-removed live worktrees [19]; and a refinery aliveness check that trusted stale pid files never respawned dead refineries [20]. These reports are evidence for failure classes, not proof that the projects remain vulnerable at later heads.

A. The Transaction Subject

The minimum authority subject is not “agent a edits repository r .” A useful abstraction is

$$A = (m, g, v, a, i, f, w, p, b),$$

where m is an immutable Mission revision, g its graph revision, v a Variant, a an Attempt, i the process incarnation, f its monotonically ordered fence, w a workspace generation, p a policy generation, and b a bounded budget reservation. A mutating request that omits one of these fields can accidentally alias a different logical run. An expired lease does not merely change a timestamp: reclaim changes i and f , so every old request is stale even if its bearer still possesses a valid-looking process handle.

The engineering subject is similarly richer than a commit hash:

$$C = H(\text{domain}, A, o_{\text{base}}, o_{\text{head}}, o_{\text{tree}}, h_{\text{patch}}, e, p),$$

where the OIDs are algorithm-tagged, e identifies admitted environment and tool subjects, and all variable-length inputs are framed. The formula is a design description, not a claim that every current constructor already carries every term. It explains the invariants: equal content with different ancestry or authority need not be the same Candidate; a moving branch is not a Candidate; and a verification result for C says nothing about C' after a rebase, merge, policy change, or gate change.

An effect has its own identity rather than borrowing the Candidate identity:

$$X = (C, E, r, o_{\text{expected}}, o_{\text{desired}}, k, f),$$

where E is the independent Evidence root, r is the protected destination, k is the logical idempotency key, and f is rechecked at the broker. The desired Git OID can be observed independently; the transport response cannot. The ledger therefore distinguishes command acceptance, local application, verification, effect intent, remote observation, and projection. Collapsing those into one “done” bit destroys the information needed for recovery.

B. Trust-Boundary Attacks

The model is intentionally hostile to ambient convenience. Repository files may request clean/smudge filters, external diff programs, merge drivers, credential helpers, hooks, submodules, alternate object stores, filesystem monitors, pagers, editors, or URL rewrites. Provider output may contain terminal escape sequences, unbounded JSON, paths outside scope, or shell names disguised as gate requests. A runner may deliver a valid request twice, reorder a stale heartbeat after reclaim, or

TABLE II
FAILURE CLASSES, UNSAFE SHORTCUTS, BINDING RULES, AND EXTERNALLY VISIBLE REFUSAL.

Failure class	Unsafe shortcut	Required binding	Honest observable outcome
Authority aliasing	PID, tmux name, or lease expiry treated as ownership	Attempt + incarnation + fence checked at every mutation	Stale principal is refused; current work remains unaffected
Repository aliasing	Parallel sessions write one checkout or common refs	Sole daemon writer and private generation per Attempt	Second writer has no repository capability
Subject drift	CI/reviewer names branch, PR, or working directory	Candidate exact base/head/tree/patch/environment; Evidence names Candidate	Rebase/merge/gate change invalidates Evidence
Ambiguous effect	Timeout classified as failure and retried as new work	Durable effect intent, one idempotency key, expected-old-OID CAS, read-back	UNKNOWN until exact applied/absent/divergent observation
False liveness	Idle terminal or missing heartbeat classified as completion/death	Lease/fence plus reconciled process and ledger observations	ALIVE, DEAD, UNKNOWN, or CONTRADICTORY
Destructive recovery	Force-remove before preserving unique bytes	Sealed preservation subject and durable cleanup/tombstone outcome	Cleanup refuses or returns UNKNOWN; recoverable bytes remain
Projection promotion	UI infers green from missing/stale state	Sequence/watermark-bound projection with no authority methods	Stale/absent/contradictory state remains visible

crash between local mutation and response. A forge may apply a write and lose the response. A projection may reconnect behind its watermark. Each is treated as a typed protocol case, not an instruction to “try again.”

Conversely, Bullet Farm does not claim to defeat an administrator who can replace the running kernel, signing roots, or admitted binaries. It does not make a same-UID path race impossible where an external tool still resolves a mutable path. The design shrinks and names the trusted computing base; it does not abolish one. Release acceptance therefore includes binary-digest admission, descriptor-relative/no-follow mutation, signed packages, provenance, and platform-specific containment rather than assuming that a source build is an installer.

III. FIVE-PLANE TRANSACTION ARCHITECTURE

Figure 1 shows the complete binding chain. The five planes are separated so that no plane can use its own assertion as independent evidence.

A. Control and Execution

The control plane owns Missions, immutable graph revisions, commands, leases, fences, policy snapshots, and durable outbox intent. A lease is insufficient without a monotonically changing incarnation fence checked by every mutation gateway [21], [22]. Reclaiming expired work produces a new authority subject; the old process cannot regain authority by waking up.

The execution plane supervises heterogeneous providers through admitted, bounded protocols. A provider receives read context and returns a typed patch proposal. It does not receive repository authority. Exactly one BulletGit daemon is the repository writer. Each Attempt uses a private clone rather than a linked worktree, eliminating shared index/common-ref control paths between Attempt principals. Hostile repository configuration, hooks, alternates, credential helpers, and environment are treated as input to be cleared or refused, not inherited ambient authority [23]. The committed policy names Landlock, seccomp, and cgroup v2 as the Linux production containment reference [24]–[26]; at the snapshot only the user and network namespace egress component exists [12], [27], so a bubblewrap-class filesystem and process sandbox for the provider [28] remains Designed.

B. Candidate and Verification

Applying a validated proposal produces an immutable Candidate identity over the provenance needed to reproduce it: base OID, head OID, tree OID, canonical patch identity, Attempt/fence, and relevant environment and policy subjects. A content hash alone is insufficient because identical bytes can arise under different authority or ancestry. A branch name is insufficient because it can move.

The verification plane reconstructs that exact Candidate in an independently clean environment. Gates come from an admitted catalog, not model-provided shell. Evidence binds the Candidate and gate IDs, executable/environment subjects, bounded output, exit status, and timing. Any subject change after verification invalidates the Evidence. A reviewer may be a different model or a deterministic tool; independence is an authority relation, not a persona label.

C. Delivery, Audit, and Ambiguity

The delivery plane alone holds forge credentials. Before an effect it durably records logical effect identity, expected-old OID, intended-new OID, Evidence root, and idempotency key. Protected integration tests the exact merge result, following the merge-queue principle that what is tested should be what becomes the branch [29], [30]. GitHub publication, when enabled, is a one-way mirror rather than a second writer: expected-old-OID CAS, read-back, divergence quarantine, and independent backups.

If transport becomes ambiguous, the command remains UNKNOWN. The broker reads the authoritative ref or remote object. An exact match can be adopted with a receipt; a mismatch is quarantined; absence may admit a retry under the same logical effect identity. HTTP success is not verification, and transport failure is not proof of non-application. This is the idempotent receiver discipline applied to Git effects [3].

The audit plane projects append-only events and exact watermarks. Portal state may be PENDING, APPLIED, VERIFIED, FAILED, or UNKNOWN; stale or contradictory streams remain visibly stale or unknown. It cannot mint authority, repair evidence, or reinterpret missing data as success.




Fig. 1. Five trust planes and the cross-plane binding chain. A timeout after effect intent produces durable UNKNOWN; read-back either adopts the exact observed effect or quarantines divergence. It never authorizes a blind retry.

D. Worked Failure and Recovery Transaction

Consider the failure class in Gas City issue 5511, where a timeout after a server-side commit led to additional workflows [16]. In the Bullet design: (1) a command and one effect identity are committed before dispatch; (2) a fenced Attempt in a private clone yields Candidate C ; (3) an independent verifier yields Evidence $E(C)$; (4) the broker records an intent to advance protected ref r from o to n ; (5) the network times out; and (6) the command becomes UNKNOWN, not failed and not complete. Read-back then has three outcomes. If $r = n$, the broker adopts the applied effect. If $r = o$, it may retry the same effect identity. If r is neither, integration is quarantined as divergence. A new runner incarnation cannot create a second logical effect because the ledger, fence, expected-old OID, and idempotency key are all checked at the mutation boundary.

The chain is the product: incarnation-fenced authority → sole repository writer → private Attempt clone → exact-OID Candidate → independently bound Evidence → durable effect intent → read-back/adoption → truthful projection. A sandbox, reviewer, or merge queue can strengthen one link but does not imply the chain.

E. Cross-Plane Invariants

The architecture is constrained by invariants rather than by agent roles. The current registry distinguishes schema controls, gateway controls, and tests so that prose cannot be the only location of a safety rule. The following ten are the compact paper view; the generated registry and closure plan remain the authoritative crosswalk [22], [31].

- 1) **One current incarnation.** At most one live incarnation/fence may hold mutation authority for an Attempt. Lease expiry alone does not grant the successor authority; the successor's durable fence does.
- 2) **One repository writer.** Provider, runner, verifier, portal, and effect processes lack direct workspace mutation capability. Only BulletGit accepts a validated operation under the exact reservation.
- 3) **Private repository generations.** Attempt generations do not use linked-worktree `.git` files or share mutable

repository control state. Prior-or-complete-next publication is preserved across crashes.

- 4) **Scoped atomic patch.** Every operation is validated before any write. Traversal, `.git`, symlink, case-fold, normalization, duplicate, special-file, body-size, and operation-count ambiguity refuse the entire batch.
- 5) **Candidate identity is immutable.** Base, head, tree, patch, and authority/provenance subjects are length-framed and domain-separated. Content identity and provenance-bound Candidate identity are not conflated.
- 6) **Evidence is non-transferable.** The verifier reconstructs an exact Candidate with admitted gates. Any change to Candidate, gate, environment, policy, or verifier subject makes prior Evidence inapplicable.
- 7) **Attestor is not broker.** The principal that produces Evidence cannot use that act to mint protected-ref authority. The effect broker checks Evidence and current control authority independently.
- 8) **Effect identity survives retry.** Intent is durable before transport. Expected-old OID, desired OID, idempotency key, and fence are read-back/adoption subjects, so retry never creates a second logical effect.
- 9) **Preservation precedes destruction.** Unique workspace bytes, journal identity, destination, and cleanup target are sealed and rechecked. Post-delete ambiguity remains an outcome, not retrospective success.
- 10) **Projection is non-authoritative.** Commands return pending before settlement. Sequence gaps, conflicting events, absent routes, and unavailable observations yield stale or unknown state; UI actions still pass through the authenticated command boundary.

These invariants deliberately retain separation among three sovereign truths. The Kernel decides whether an operation is authorized. BulletGit decides what repository subject exists. The verifier decides whether Evidence for that subject satisfies admitted gates. The broker decides whether an authorized, verified subject may be proposed to the forge, while the forge remains sovereign over its protected ref. No single database row replaces those independent observations.

F. Evolutionary Control as Design

The intended product is a multi-agent system in which several Variants compete inside one Mission under the invariants above. Everything in this subsection is **Designed**: the normative model is a reviewed document, the committed policy fixes `evolutionary_authority=false`, and the closure plan gates the program on a connected transaction [22], [24], [32]. “Evolutionary” means bounded search over immutable, evidence-bearing software-change lineages; not autonomous credential use, self-reported fitness, unbounded agent debate, or online self-modification.

The genome is an immutable, content-addressed `TeamRecipe`: typed roles and contracts, communication edges and context policy, certified provider/profile choices, bounded budgets, and deterministic stopping, escalation, dissent, fusion, and fallback rules. A recipe cannot contain or evolve credentials, authority claims, safety rules, risk classes, protected paths, evidence floors, hidden evaluators, integration policy, effect permission, or hard budget ceilings, and a new recipe inherits no live session, credential, workspace, proof, or cache. A deterministic feasibility shield (authority, containment, certification, role conflicts, quota, budget, verifier capacity, task policy) runs before any objective, so infeasible recipes consume no budget and no score offsets a hard constraint. Eligible results enter a bounded quality-diversity archive in the MAP-Elites family [33], [34]: human-versioned, receipt-observable descriptors define cells, and each cell keeps a Pareto set over independently measured correctness, robustness, cost, latency, intervention, and observation survival with a content-ID tie-break; there is no global scalar reward.

Recipes advance only through completed, comparable rungs B0–B6 in the successive-halving discipline of staged budgets [35], [36]: deterministic checks; a small matched visible block against the incumbent; a larger block under the same total compute; a sealed confirmation under independent oracle custody; no-effect shadow routing; a bounded canary with automatic rollback; and an eligible routing arm with an incumbent traffic reserve. Raw holdout failures never enter author context. Every task class keeps a strong single-agent incumbent; agent count, transcript volume, confidence, and activity are never fitness dimensions; a sparse island model copies only immutable recipe identity, re-evaluated at the destination [37]; and a monitored bound breach removes a recipe and falls back to the incumbent. V1 records routing and outcome provenance and does not update weights, prompts, policies, or privileges online. Contextual routing learners [38] are post-V1 and require holdouts with unbiased offline evaluation [39], rollback, drift detection, and an explainable signed decision record. None of these mechanisms has a receipt; they are absent by design rather than incomplete.

G. BulletGit as the Repository Trust Boundary

BulletGit is not intended to replace Git’s object model or pack protocol. It wraps ordinary Git with a capability and evidence boundary. A bare mirror can amortize network/object transfer, but an Attempt clone dissociates before

use so no surviving alternates file makes mirror garbage collection part of workspace correctness. The clone loses its origin and ambient authentication. Git child processes receive a scrubbed environment and explicit configuration that disables inherited global/system config, paging, hooks, credentials, external attributes, signing, editors, and prompts. Repository-scoped commands first inspect local configuration with includes disabled; command-bearing or truth-redirecting values fail closed [23].

This distinction matters because a repository is active input. A clean filter can execute during checkout; `textconv` can alter what a diff appears to show; an external merge driver can create bytes outside the intended patch; URL rewriting or a remote helper can redirect network authority; a credential helper can make a private clone push-capable; alternates can make object availability depend on another repository; and hooks can turn an apparently read-only Git verb into execution. A provider sandbox does not automatically neutralize those semantics, especially when the provider itself runs Git.

Patch application validates the complete batch before mutation. Paths are canonical normalized relative bytes with no parent traversal, `.git` component, NUL, backslash ambiguity, alternate-data-stream syntax, trailing dot/space, symlink traversal, or folded collision. Operation and body counts are bounded. Checkpoints use a temporary index rather than the live workspace index. A Candidate is minted only after Git reports the expected tree/head relationship and the journal durably identifies the generation.

Preservation is part of the protocol rather than cleanup hygiene. A sealed receipt binds source workspace, recoverable artifact, destination, journal, Candidate/Attempt subjects, and exact cleanup target. Immediately before deletion the daemon re-verifies the receipt and the filesystem object. A synced tombstone records the settled cleanup outcome. If failure occurs after deletion but before confirmation, the result is typed `UNKNOWN`; the implementation does not reconstruct certainty from absence.

The current production constructor is intentionally fail-closed: positive online final-authority settlement is unavailable, so the public mutation path refuses before repository I/O. That refusal is useful component evidence. It is not evidence that the designed positive path works. Enabling it requires an immutable signed shared-wire subject, locally verified grants, current Kernel reservation/fence read-back, short operation permits, and durable settlement of the same transaction.

H. Jeryu Source Authority and Blocked Topology Proposals

The frozen V1 topology distributes source Jeryu-only from immutable signed tags. It assumes neither `git.neverhuman.org` nor any GitHub namespace or public mirror, and the family manifest admits no authenticated Jeryu URL. The only permitted local Jeryu family is `/home/ubuntu/jain-split/jeryu-split`; `/home/ubuntu/jeryu-split` must never be recreated.

The named `git.neverhuman.org` deployment and `github.com/neverhuman/jeryu` and `github.com/neverhuman/jankurai` front doors are

unratified, operator-blocked proposals, not current topology or authority. A later operator choice may treat GitHub as a public index or backup and host selected open-source families on independent Jeryu Git; neither name is source authority today. If a public mirror is ratified, it remains a separately configured effect adapter: expected-old-OID CAS, read-back, divergence quarantine, and independently proved backups. A mirror is not a backup.

Jeryu’s local hub provides branded onboarding, the family manifest, and pinned component mapping while component repositories remain independent. Jankurai is different: it is an existing monorepo, but no public front door is assumed. Before any external-history or mirror decision, the transaction rule is preservation first—complete ref and object bundles, ancestry map, reconciliation RFC, rehearsal, and explicit approval. A force-push would discard the very evidence required to decide which history should survive. The local-family rule is frozen. Endpoint and mirror proposals remain **Designed** and operator-blocked unless dated ratification precedes TLS, authentication, protection, push/check/read-back, backup/restore, deployment, and mirror receipts.

IV. IMPLEMENTED EVIDENCE AND CONNECTED GAPS

Table III separates current component receipts from the connected claims they do not establish. Repository subjects are drawn from the shared evidence snapshot: Hub 3039878371b6, Kernel ca380bc4ffd4, BulletGit fb715b25c3b2, and Portal 95108e37d643. A later source commit invalidates mapped receipts until replay; publication-only descendants are admitted only when preflight proves that evidence-critical inputs did not change.

The table is intentionally asymmetric. Numerous component tests can coexist with zero connected-product proof. just demo is a deterministic demonstration, not a live transaction. The doctor result is **BLOCKED** because the checked family lock is intentionally unsupported and stale. The release inventory is a truthful negative receipt, not a score.

At the Kernel subject the live-admission boundary is mechanized on both sides and closed. The Kernel policy loader mirrors the Hub’s v1alpha2 rule (live admission only at policy generation ≥ 2 with an active provider-runner key; commit 0d848f6); the nightly lane’s real-binary mode refuses to start without an operator-ratified policy path (b4735da); and under the committed generation-1 policy every provider is refused at `POLICY_LIVE_ADMISSION_DISABLED` before any key read, probe, or spawn, with default nightly halves targeting a marker script so that any spawn fails the lane. Ratification is an operator act recorded in the family ledger. No provider has a live receipt [12], [22].

A. Evidence Lifecycle

An admissible receipt answers five questions: *who observed what exact subject, using which admitted mechanism, under which boundary, with what outcome?* Test counts alone answer none of them. The paper snapshot consequently stores repository commit and tree, clean-state fact, tool versions, complete command, exit code, semantic outcome, and a sentence stating

what the result does not prove. It also hashes the family manifest, family lock, release truth, invariant registry, formal locks, policy, and competitor snapshot. These inputs are deliberately outside the PDF hash loop.

Evidence moves monotonically only within one immutable subject. A new commit, tree, policy generation, executable, gate, environment, or dependency lock does not make the old receipt false; it makes it inapplicable to the new subject. Replaying a command in a dirty tree is useful diagnosis, but its result cannot be committed as frozen evidence because the tested bytes are not named by the repository commit. Similarly, a generated release report is a projection of registered receipts and blockers. It cannot satisfy one of its own rows.

Component receipts are still valuable. Atomic Kernel tests can falsify a lease or command implementation. Hostile Git fixtures can falsify configuration or path handling. Browser tests can falsify projection behavior. Bounded model checks can falsify protocol invariants under their state bounds. The discipline is to preserve those results without performing an invalid conjunction: “Kernel passed” $\wedge$ “BulletGit passed” $\wedge$ “Portal passed” does not imply that a signed command traveled through all of them with one Candidate and one observed effect.

B. Formal Models and Their Boundary

Exactly two bounded TLA+ models are pinned and checked with TLC [40], [41]; the lock fixes the TLC release by digest, the Java major, single-worker execution, and the module, configuration, and state counts. At the snapshot the model lane is **PASS** for 2 models. *LeaseFence* models two runners, one lease, a monotonically increasing fence, a database clock, scope revisions behind a mutation barrier, freeze, and restore epochs (MaxFence=3, MaxTime=4, MaxScope=3, MaxOps=2, MaxEpoch=2, MaxFreeze=3); it checks exactly `TypeOK`, `BarrierHasNoInFlightApply`, and `RestoredAuthorityNeedsFreshAcquire` over 529,851 generated and 128,913 distinct states at depth 24. Stale-fence, stale-epoch, expired, frozen, and barrier applies are refused by the authorization guard and counted, not asserted as an invariant, and fence monotonicity holds by construction (the fence changes only on acquisition, by one), not by check. *EffectCheck* models one effect and one dependent check against a remote that may be untouched, desired, or written by a third party (MaxDispatches=2, MaxEpoch=2); it checks exactly `TypeOK`, `AtMostOneLogicalEffect`, `VerifiedEffectWasReadBack`, `CheckRequiresDurableProof`, and `ThirdPartyStateIsNeverAdopted` over 1,585 generated and 378 distinct states at depth 18. `AtMostOneLogicalEffect` is enforced by assignment rather than by counting physical writes, so the dispatch guard carries the real protection, and `NoDispatchAfterStop` is defined but not listed in the checked configuration.

Both configurations set `CHECK_DEADLOCK FALSE`; no liveness or fairness property is checked. Wound-wait [42], multi-resource scope expansion, two Attempts on one Variant,

TABLE III
IMPLEMENTED NOW VERSUS NOT YET PROVED AT SNAPSHOT 20260825-STAGE1-R3.

Boundary	Maturity	Implemented evidence	Not established
Wire and admission	Component-proved	Canonical signed DTOs, bounded IPC, executable-digest and launch-grant components	Immutable shared-wire publication and connected issuer/consumer settlement
Control authority	Component-proved	DB-clock leases, fences, atomic command/outbox, authenticated ingress, quarantined signed lease-transport, proposal-to-checkpoint binding	Durable Mutation reservation/settlement and one connected Kernel-Runner-BulletGit issuer path
Provider boundary	Component-proved	Four offline message subsets, policy-first refusal, bounded process and Linux egress components	Any live provider receipt; policy remains live-disabled in the frozen evidence subject
BulletGit	Component-proved	Fail-closed production gateway, private-clone/preservation primitives, cleanup UNKNOWN, strict local manifests	Positive online sole-writer transaction; public clone remains <code>AUTHORITY_CONTRACT_UNAVAILABLE</code>
Verifier	Component-proved	One admitted fixture gate, exact-subject Evidence, bounded transport failure handling	Production reconstruction source, executable digest for every gate, and multi-gate aggregation
Effects	Component-proved	Command/effect ambiguity model and local bare-forge components	Authenticated Jeryu/GitHub expected-old-OID push and read-back receipts
Portal/audit	Component-proved	Generated contracts, authenticated command reconciliation, watermark-bound projections, explicit absent-surface UNKNOWN	Packaged runtime and complete Candidate/Evidence/Effect ledger subjects
Security audit	Component-proved	Deterministic score 58 (raw 58), 10 caps, 29 hard and 15 soft findings	V1 floor ≥ 90 , zero caps/hard findings, and a portable pinned hosted artifact
Connected product	Designed	Normative five-plane transaction and release contracts	Transaction-proved or Live/release-proved evidence; 26 blocked of 26 evaluated release gates

TABLE IV
STAGE-1 COMMAND FAMILIES, RECORDED OUTCOMES, AND INTERPRETATION. FULL COMMANDS AND BOUNDARIES ARE IN `EVIDENCE.JSON`.

Command family	Snapshot outcome	What it establishes	What it cannot establish
<code>bullet-family doctor -json</code>	BLOCKED	Fail-closed install/lock diagnosis	A supported installed family
<code>bullet-family check release -json</code>	BLOCKED	26 named blockers among 26 evaluated gates	Any positive gate or percentage score
<code>Hub just fast</code>	PASS	Mapped Hub component checks at the subject	Sibling, transaction, live, or release truth
<code>Hub just demo</code>	PASS	Deterministic multi-process demonstration semantics	<code>TRANSACTION_PROOF</code> ; provider/forge/release evidence
Hub contract lane	PASS	Wire/schema generation and drift checks	Immutable publication or consumer settlement
Invariant registry	PASS	Machine crosswalk and enforced-entry validation	Implementation of rows whose lifecycle is planned
Local audit lane	PASS component; release blocked	Score 58/raw 58, 10 caps, 29 hard and 15 soft findings	V1 quality floor or portable hosted execution
Repository required lanes	PASS	Mapped committed components for Kernel, BulletGit, and Portal	Cross-repository composition
Formal model lane	PASS	2 bounded models match pinned sources/tools/counts	Liveness or implementation correctness
Paper build/preflight	deterministic publication check	Shared values, citations, layout, metadata, embedded fonts, page size/count, reproducible PDF bytes	Product correctness or release authority

Git, provider protocols, the forge, OS containment, and multi-tenant attackers are not modelled. The JSON traces beside the models are executable fixtures replayed by Rust tests, not model-check results. Model success therefore remains Component-proved, never Transaction-proved.

C. Open Connected Product Boundary

The authoritative closure plan, not this paper, controls V1 [22]. The remaining predecessor chain includes an immutable `bullet-wire` tag that current consumers may honestly resolve, a Kernel reservation that repeats lease checks inside one write, production BulletGit verification of that permit, connected Candidate/Evidence settlement, authenticated protected integration, provider conformance, packages, installer, backup, signing, provenance, and live/release receipts. The new workplan is only an opportunity backlog and explicitly loses to that authority [43]. Most importantly, the absence of an end-to-end `TRANSACTION_PROOF` is not hidden behind component totals.

V. PINNED SYSTEMS AND CAPABILITY MATRICES

The comparison asks two different questions. Operational capabilities concern what work a system helps users perform. Authority/integrity properties concern what exact subject and principal the public contract binds. Combining them would incorrectly make a rich orchestrator look unsafe merely because it is featureful, or make a strict component look operationally complete merely because it refuses unsafe work.

A. Deep Case Studies

Gas Town v1.2.1. Gas Town provides a legible multi-agent town: role-specific agents, Beads-backed work, tmux sessions, hook worktrees, patrol and recovery machinery, and a Refinery merge queue [5], [44]. Its strength is operations: many agents and providers can be supervised using durable work records. Its public contract uses worktrees and orchestration roles rather than a separate repository-writer principal. Because linked worktrees have their own `HEAD` and index but share the common Git repository and most refs, they are not equivalent to private repository principals [13].

Gas City v1.4.1. Gas City generalizes the factory with Agents, Beads, Formulas, Rigs, Packs, Events, multiple run-times, retry/fan-out/drain controls, and release artifacts [6], [45]. The richer runtime and workflow vocabulary should not be minimized. The issue audit nevertheless shows why session durability and exactly-once effects are distinct: a shared rig root, a non-atomic task claim, and retry after ambiguous creation all cross different trust boundaries [15]–[17].

DeepSeek Harness rc.2. DeepSeek Harness is a Cordis plugin kernel: model, tools, sandbox, storage, session, loop, scheduling, and UI are composed as plugins, with model-visible context reconstructed from an append-only session log [46]. It is not accurately described as categorically single-agent. Its experimental Agent Teams seam adds a roster, task board, mailbox, and continuable child agents. The same-process team documentation also warns that checkout scopes are advisory rather than locks [47]. Its primary contribution is composable agent execution, not protected Git integration.

Omnigent v0.10.0. Omnigent wraps several vendor harnesses behind a uniform session/runtime layer. It documents Linux bubblewrap and macOS seatbelt sandboxes plus an L7 proxy that keeps credentials outside the sandbox and can transform network requests [48], [49]. Contextual policy verdicts are ALLOW, DENY, and ASK; transformation belongs to the L7 proxy, not the policy verdict [50]. Omnigent Polly adds parallel worktrees and cross-vendor implementation/review, which is a meaningful operational feature [51], [52]. The wrapped harness still performs repository actions, so approval attribution and Candidate binding remain separate questions.

B. What the Architecture Should Borrow

The comparison is not a feature-stripping exercise. Gas Town demonstrates the operator value of named roles, inspectable work records, patrol, recovery, and a dedicated landing function. Gas City makes runtimes and workflow structures more explicit and adds release-operational depth. A transaction processor that cannot explain where work is, recover a session, limit dispatch, or expose a usable queue will be safer only by being unusable. Bullet’s control plane should compose these operational concepts while ensuring that a role name never substitutes for an authority subject.

DeepSeek Harness shows that provider, tools, storage, sandbox, loop, and UI can be composable plugins without turning every integration into a special case. Its append-only model-visible session log is especially relevant: if context cannot be reconstructed from durable events, recovery silently changes the agent’s inputs. Bullet can consume a plugin kernel behind the provider boundary as long as plugin composition cannot acquire repository or effect authority. Experimental Teams also demonstrates that subagent coordination and a task board can be first-class protocol objects. The remaining difference is enforcement: advisory checkout scopes are useful coordination, while a transaction gateway must refuse overlapping mutation.

Omnigent demonstrates mature defense in depth at a different layer. Native OS sandboxing, a credential-isolating L7 proxy, contextual approval, shared sessions, and cross-vendor review are desirable even when sole-writer and Candidate

bindings exist. Bullet should not attempt to reimplement every agent UX. It can wrap or accept proposals from those systems, provided the proposal is bounded untrusted data and all repository/effect credentials remain across the gateway. Conversely, an L7 proxy can transform or authenticate network traffic without becoming the authority that says a Candidate passed its gates.

These examples suggest a clean interoperability rule: protocols such as ACP and MCP may carry sessions, context, and tool calls, but they do not mint Bullet capabilities. A provider adapter may translate a typed proposal. A review adapter may return commentary. An MCP tool may expose read-only repository facts. None may cross from “the model requested” to “the repository is authorized” without the same signed control subject, current fence, admitted operation, and durable settlement used by native components.

C. Classification Procedure

For each cell, the review first fixed the release/commit and source set; then looked for an explicit contract; then assigned the least affirmative supported term. A feature shown only in an optional pack, plugin, hosted configuration, or experimental seam is **Partial/configuration-dependent** unless the row itself is about that seam. **Not documented** records a search result, not an implementation claim. **Unknown** is used when the public description is too ambiguous even to say that the property is undocumented. **N/A** is reserved for products whose public role does not include the property—for example, a plugin kernel that does not claim to operate a protected forge.

The Bullet design and Bullet Stage-1 rows are separate. Otherwise a design document would be compared against released competitors as if it were running code. The current row is **Partial/configuration-dependent** whenever some exact component exists but the positive connected path does not. A fail-closed unavailable production constructor is evidence that unauthorized writes are refused; it is not evidence of an operational repository writer. This conservative classification also prevents the matrices from becoming marketing scores: the terms are not ordered numerical values and should not be summed.

D. Why Familiar Mechanisms Do Not Automatically Compose

One objection is that every part of the chain already exists: worktrees isolate files, containers isolate processes, CI verifies commits, merge queues protect branches, and idempotency keys prevent duplicate requests [53]. The premises are mostly correct; the conclusion does not follow unless the subjects and principals at their boundaries are the same.

A worktree is not a writer identity. It gives a convenient per-worktree HEAD and index, yet common repository metadata and most refs remain shared [13]. An agent that can invoke arbitrary Git can still name refs, helpers, remotes, hooks, config, objects, and sibling worktrees outside the intended edit surface. A private clone with authentication removed is not magical isolation either, but it makes the repository principal a separable object that the sole writer can validate and destroy.

A sandbox is not a repository transaction. A container, bubblewrap [28], seatbelt profile, or remote VM can restrict filesystem and network access. It does not decide whether the bytes being written are an admitted patch for the current fence, whether tests name the resulting tree, or whether a protected push applied after timeout. Sandboxing is a necessary defense against provider and repository content; it complements rather than replaces Candidate and effect identity.

A durable session is not durable engineering state. Transcripts, mailboxes, task boards, tmux sessions, and event logs make reasoning resumable. They may prove what an agent saw or said. They do not prove that the repository contains a particular tree, that a gate ran against it, or that the forge accepted it. Bullet should retain session history for cognition while deriving repository truth from Git and verification truth from Evidence.

A reviewer is not automatically an independent attester. A second model can catch semantic mistakes and cross-vendor review can reduce correlated failure. Independence still fails if both reviewers share a mutable checkout, if the author can rewrite the test oracle, if the review result names only a branch, or if the reviewer credential can execute the protected effect. The authority requirement is reconstructible exact-subject verification and attester/broker separation; model diversity is an additional quality control.

A merge queue is not effect reconciliation. Testing the exact merge result is a central integrity property [29], [30]. The queue still needs a caller identity, current authority, an Evidence subject, expected-old ref, and a rule for lost responses. A queue item reported as failed after its merge applied can be duplicated by an upstream retry unless logical effect identity survives transport ambiguity.

An idempotency key is not useful unless the receiver binds it. A random client token stored only in process memory disappears on restart. A receiver that records the token after rather than atomically with its effect leaves a crash window. A token reused for a different desired OID launders a conflict. The effect record must bind the logical command, Candidate, Evidence, protected ref, expected and desired OIDs, current fence, and observed outcome, durably before and after transport [3].

Bullet Farm’s novelty claim is therefore compositional and falsifiable, not that it invented locks, hashes, queues, or idempotence. If a smaller public contract can demonstrate the same bindings under the fault corpus, Bullet’s extra machinery is unnecessary. Until then, comparing one isolated mechanism with the entire chain is a category error.

The result is not that Bullet Farm is universally superior. It has less live product evidence than every mature system in this table. The defensible result is that the reviewed pinned public contracts do not document the same end-to-end bindings. Whether those bindings justify their cost is precisely the Stage-2 benchmark question.

VI. ISSUE-TO-INVARIANT AUDIT

Table VII records state and affected/reported versions as observed on 2026-08-25. “Pinned corroboration” distinguishes

issue testimony from code/release material at the selected subject. An issue being closed does not erase the failure class; an open issue does not prove current pinned code is affected.

The audit corrects four tempting misreadings. First, issue 742 is about a settings path; it supports no respawn-count claim. Second, DeepSeek Teams makes categorical “single-agent” language false. Third, Omnigent policy returns ALLOW/DENY/ASK; proxy transformations are a different layer. Fourth, linked worktrees do not share one index: each linked worktree has its own index and HEAD, while common repository metadata and most refs are shared [13].

VII. EVALUATION, FALSIFICATION, AND ARTIFACT

A. Stage-1 Evaluation Protocol

Stage 1 is architecture and component assurance. From one dated evidence snapshot, the artifact records repository commit/tree/clean facts, toolchain versions, full commands, doctor and release-gate outcomes, component and formal receipts, external pins, and hashes of evidence-critical inputs. Both PDFs consume generated macros from that JSON. The strict preflight verifies the schema, subjects, trees, cleanliness, input hashes, citations, references, layout warnings, metadata, embedded fonts, letter page size, page counts, and byte-identical rebuilds at one source epoch. It deliberately fails while any canonical family checkout is dirty.

The reproducible command surface from the Hub is:

```
just paper just paper-check
cargo run -locked -quiet -bin
bullet-family - doctor -json
cargo run -locked -quiet -bin
bullet-family - check release -json
just fast just demo
```

Repository-mapped required/contract/formal commands and their complete observed outcomes are enumerated in `evidence.json`; summary macros do not replace those command records. The Hub evidence subject excludes publication-only descendants to avoid an impossible self-hash loop. Preflight permits that narrow descendant only when every changed path is in the paper/tooling set and all hashed evidence inputs remain identical.

B. Fault-Oriented Acceptance Corpus

A happy-path patch is a weak discriminator: nearly every reviewed system can edit a file and run a test. The connected Bullet acceptance corpus must inject faults at every state boundary and retain the resulting ledger, Git, Evidence, effect, and observation subjects. At minimum it includes:

- 1) reclaim an Attempt while its old provider and runner remain alive, then deliver old mutation, heartbeat, verification, and effect messages;
- 2) crash before and after patch journal sync, Candidate publication, Evidence settlement, effect-intent sync, remote application, response receipt, read-back, preservation, deletion, and tombstone sync;
- 3) rebase or merge between verification and integration, change a gate executable or policy generation, and attempt to reuse stale Evidence;

TABLE V
OPERATIONAL CAPABILITIES IN PINNED PUBLIC CONTRACTS. THESE CELLS DO NOT IMPLY AUTHORITY SAFETY.

System	Orchestration	Provider composition	Isolation	Durable sessions	Review	Merge queues
Bullet Farm Stage 1 [4], [22]	Documented	Partial/configuration-dependent	Partial/configuration-dependent	Documented	Partial/configuration-dependent	Partial/configuration-dependent
Gas Town v1.2.1 [5], [44]	Documented	Documented	Partial/configuration-dependent	Documented	Partial/configuration-dependent	Documented
Gas City v1.4.1 [6], [45]	Documented	Documented	Partial/configuration-dependent	Documented	Partial/configuration-dependent	Documented
DeepSeek Harness rc.2 [46], [47]	Documented	Documented	Partial/configuration-dependent	Documented	Documented	Not documented
Omnigent v0.10.0 [48], [51]	Documented	Documented	Documented	Documented	Documented	Partial/configuration-dependent

TABLE VI
AUTHORITY AND INTEGRITY PROPERTIES. NOT DOCUMENTED MEANS ONLY THAT THE REVIEWED PINNED PUBLIC CONTRACT DOES NOT STATE THE BINDING.

System	Writer identity	Incarnation fence	Exact verification subject	Effect idempotency/read-back	Protected integration	Truthful uncertainty
Bullet Farm design [4]	Documented	Documented	Documented	Documented	Documented	Documented
Bullet Farm Stage 1 [12], [22]	Partial/configuration-dependent	Partial/configuration-dependent	Partial/configuration-dependent	Partial/configuration-dependent	Partial/configuration-dependent	Documented
Gas Town v1.2.1 [5], [44]	Partial/configuration-dependent	Not documented	Not documented	Not documented	Partial/configuration-dependent	Partial/configuration-dependent
Gas City v1.4.1 [6], [45]	Partial/configuration-dependent	Not documented	Not documented	Partial/configuration-dependent	Partial/configuration-dependent	Partial/configuration-dependent
DeepSeek Harness rc.2 [46], [47]	Not documented	Not documented	Not documented	N/A	N/A	Partial/configuration-dependent
Omnigent v0.10.0 [48], [50]	Partial/configuration-dependent	Not documented	Not documented	Not documented	Partial/configuration-dependent	Partial/configuration-dependent

TABLE VII
PRIMARY ISSUE EVIDENCE MAPPED TO BULLET INVARIANTS AND CURRENT MATURITY.

Project / ID / state / version	Pinned-code or release corroboration	Failure class	Mapped Bullet invariant	Current maturity	Residual risk
Gas Town #4094; open; reported 1.1.0 [14]	v1.2.1 retains town/hook worktree architecture; issue is not treated as a pin-wide verdict	Repository/config aliasing	Sole writer; hostile Git/config refusal; private Attempt clone	Component-proved	Connected BulletGit authority and same-UID path races remain open.
Gas Town #742; closed; reported before v1.2.1 [54]	v1.2.1 was inspected; no “687 respawns” figure is present or claimed	Configuration placement	Sealed admitted HOME/config and exact executable subject	Component-proved	Live provider configuration and every real-provider receipt remain absent.
Gas Town #3584; closed; version not stated [18]	v1.2.1 README keeps the Problems-view row “Zombie: Dead tmux session” as the liveness signal [44]; issue records a mid-task death that left work hooked without notification	Inferred liveness (silent death)	Lease/fence heartbeat with runner self-fence; hooked work never completes by silence	Component-proved	DB-clock leases and self-fence exist; the four-valued reconciled observation is designed.
Gas Town #4737; open; reported 1.1.0 [19]	v1.2.1 retains patrol and cleanup molecules; issue records a liveness check that classified every dog dead before forced worktree removal	Inferred liveness (false death) and destructive recovery	Destructive eligibility only from exact incarnation identity; preservation before deletion	Component-proved	Preservation and tombstone are local components; no connected cleanup settlement.
Gas Town #4091; open; version not stated [20]	v1.2.1 retains daemon aliveness checks; issue records stale pid files trusted as liveness	Inferred liveness (stale identity)	Incarnation fence and ledger observation, never PID or session name	Component-proved	Fence is a component; the reconciled ALIVE/DEAD/UNKNOWN/CONTRADICTION observation is designed.
Gas City #1181; open; version not specified [15]	v1.4.1 documents rigs/packs; issue records parallel agents resolving to one rig path	Shared writable identity	Private per-Attempt clone; one writer principal	Component-proved	No connected positive BulletGit transaction.
Gas City #5511; open; affected v1.4.1 [16]	Exact selected release is the reported affected version; issue records post-commit timeout and duplicate workflows	Ambiguous effect retry	Durable effect identity; expected-old-OID CAS; read-back/adoption	Component-proved	No authenticated live forge receipt or transaction proof.
Gas City #1052; closed via #2985; affected 0.15.2 [17]	v1.4.1 is later than the reported fix; retained as historical class evidence only	Non-atomic claim	Atomic lease/fence/Attempt command transition	Component-proved	Signed online transport is not connected end to end.
Omnigent #1694; open; affected 0.2.0 [52]	v0.10.0/Polly materials document parallel worktrees and cross-vendor review	Session/config ambiguity	Durable Attempt state; incarnation fence; preservation before cleanup	Component-proved	Provider and reviewer identities are not yet live-bound in Bullet.
Omnigent #725; open; version not stated [55]	v0.10.0 scope was inspected; issue concerns external/native edits in a non-Git workspace	Incomplete change observation	Candidate manifest and honest UNKNOWN for unrepresented change state	Component-proved	Non-Git workspaces are outside Bullet V1; Git Candidate path is not connected.
Omnigent PR #4318; merged; included in v0.10.0 [8], [56]	Release notes identify the approval-attribution rollback as a breaking change	Credential/approval attribution	Separate requester, approver, broker, and effect identities	Designed	Operation-specific components exist, but no live owner-credential effect proves separation.

- 4) inject hostile Git config, hooks, filters, alternates, attributes, credentials, helpers, submodules, symlinks, path normalization/case collisions, large frames, invalid Unicode, terminal escapes, and protocol contamination;
- 5) drop, duplicate, reorder, truncate, and delay every JSON-RPC and command response, including a server-applied forge write whose client response is lost;
- 6) replace an admitted pathname after validation, race cleanup targets, exhaust descriptor-relative traversal limits, and exercise each supported or explicitly refused operating system; and
- 7) disconnect/reconnect the projection across sequence

gaps, conflicting events, missing routes, and restore epochs, verifying that no display becomes authoritative.

The acceptance criterion is not merely “no crash.” Every injected boundary must end in one of a small set of independently inspectable outcomes: exact prior state, exact complete next state, durable failed state with preserved subject, or truthful UNKNOWN/quarantine requiring reconciliation. No outcome may contain two current fences, two logical effects, Evidence for the wrong Candidate, an unpreserved unique workspace, or a projection greener than its source.

Formal and implementation traces should cross-check each other without being confused. The two TLA+ models can gen-

erate small safety traces for lease/fence and effect ambiguity. Rust tests can replay those traces against component state machines. The connected corpus must then repeat equivalent faults across real processes, filesystem boundaries, and a controlled forge. Agreement is stronger evidence than any layer alone, but it remains one offline Transaction-proved subject until live admission repeats the same binding.

C. Stage-2 Matched Benchmark

Stage 2 is prohibited until a signed connected TRANSACTION_PROOF exists. It will freeze identical repositories and issue tasks; pin every system, model, tool, policy, budget, and retry rule; run enough trials to expose tail failures; and retain exact final trees, tests, human adjudication, wall-clock time, token/cost measures, collisions, duplicated effects, ambiguous outcomes, recovery actions, and data loss. Operational success and authority violations will be separate dependent variables. Systems unable to emit a native receipt will be wrapped by an external observer without pretending that wrapper is a product feature.

Only that matched corpus can support “faster,” “cheaper,” “safer,” or “more successful.” This paper makes none of those claims. Public superiority language remains deferred with the benchmark.

Matched evaluation must also account for refusal. A system that declines every mutation has no integrity failures and no product utility; a system that always continues may look fast by discarding ambiguity. Results therefore report task utility, safe refusal rate, recovery latency, operator interventions, and unresolved unknowns together. Storage overhead from private clones, verifier reconstruction time, merge-queue dwell, and read-back traffic are first-class costs rather than hidden implementation details.

D. Falsification Conditions

The architecture would be weakened or rejected by any of the following:

- a stale or reclaimed incarnation mutates a repository or remote effect;
- two Attempts share a repository principal or common mutable control path;
- Evidence remains valid after any Candidate subject changes;
- timeout/restart can create two logical effects or erase an applied one;
- cleanup destroys the only recoverable bytes before preservation settles;
- Portal or prose can promote UNKNOWN to verified;
- the same guarantees are achieved by a substantially smaller public contract under the matched fault corpus; or
- bookkeeping/latency cost dominates benefit under realistic workloads.

The paper itself is falsified by a source row that does not support its cell, an unrecorded dirty-tree result, a snapshot mismatch, a non-reproducible PDF, or a claim above its maturity. Those are publication failures even if the Rust components remain correct.

E. Limitations and Artifact Statement

This is a design and component inventory from one date. Public documents can omit private controls; issues are selected counterexamples rather than incidence statistics; proprietary systems are black boxes; and open-source heads move after their pins. The V1 threat model does not solve multi-tenant remote execution, a malicious host kernel, semantic correctness of model output, non-Git workspaces, or every same-UID race. Private clones cost storage and setup time. Fencing, exact subjects, and read-back add state and latency. No connected-product measurement yet shows that the cost is worthwhile.

The artifact comprises this source, the executive brief, bibliography, architecture vector, evidence JSON and schema, generated macros, deterministic build/preflight scripts, and the non-authoritative workplan. The frozen V1 topology is Jeryu-only from immutable signed tags, with no authenticated URL or public mirror selected. The only permitted local Jeryu family is `/home/ubuntu/jain-split/jeryu-split`. The proposed `git.neverhuman.org` deployment and GitHub discovery mirrors remain operator-blocked and unratified. Jankurai remains its existing monorepo and requires preservation-first review before any mirror decision. None of those topology proposals is a live receipt [43].

VIII. CONCLUSION

Powerful agents increase the need for narrow, inspectable authority. Bullet Farm’s proposed contribution is neither a stronger persona nor another agent loop. It is the composition of familiar transaction mechanisms over exact Git subjects: fencing at mutation, one repository writer, private Attempts, immutable Candidates, independent Evidence, durable effect intent, read-back/adoption after ambiguity, protected integration, and truthful uncertainty across five planes.

The current evidence boundary is equally important. The family has meaningful component receipts and bounded models, while the connected transaction, every live provider/forge proof, and the release remain absent. The reviewed pinned public contracts do not document the same end-to-end bindings, but that is a contract comparison rather than measured superiority. Stage 1 publishes the architecture and its falsifiable gaps. Stage 2 begins only after the exact connected transaction exists.

APPENDIX A

DATED LANDSCAPE BEYOND THE DEEP STUDIES

Table VIII is an observation index, not a capability score. Open-source entries name the exact inspected commit; proprietary entries name official documents observed 2026-08-25 and remain black-box limited.

Release/commit pins are observation subjects, not endorsements and not claims about later versions. For proprietary systems, only the cited official surface is classified; undocumented properties remain Unknown.

TABLE VIII
LANDSCAPE SNAPSHOT OBSERVED 2026-08-25. FULL 40-HEX PINS ARE IN THE BIBLIOGRAPHY AND EVIDENCE MANIFEST.

System	Pinned subject or official observation	Public-contract scope and limitation
Codex	Official CLI/cloud docs [9], [10]	Local CLI and hosted delegated tasks are documented; public docs do not establish Bullet's connected authority chain. Proprietary/hosted behavior is black-box limited.
Claude Code	Official docs observed 2026-08-25 [57]	Agentic terminal/editor workflow; proprietary behavior and service controls are black-box limited.
Cursor	Official agent docs observed 2026-08-25 [58]	IDE/agent workflow; proprietary internals are black-box limited.
GitHub Copilot	Official coding-agent docs observed 2026-08-25 [59]	Hosted issue-to-PR workflow integrated with GitHub; service internals are black-box limited.
Jules	Official Google docs observed 2026-08-25 [60]	Hosted asynchronous coding tasks; service internals are black-box limited.
Devin	Official docs observed 2026-08-25 [61]	Hosted coding-agent sessions; service internals are black-box limited.
Kiro	Official docs observed 2026-08-25 [62]	IDE/spec-driven agent workflow; service internals are black-box limited.
Gemini CLI	commit 812f7a2bcf20 [63]	Open-source terminal agent; Apache-2.0.
OpenHands / Agent Canvas	commit 150e76046db0 [64], [65]	Open agent platform and control surface; MIT. Agent Canvas documents isolated worktree operation and ACP composition.
SWE-agent	commit 3ea751c087f3 [66]	Research issue-solving harness; MIT.
mini-SWE-agent	commit 25941c89cfeb [67]	Minimal research harness; MIT.
Aider	commit 5dc9490bb35f [68]	Open terminal pair-programmer; Apache-2.0.
Cline	commit 4f5f2384076e [69]	Open editor agent; Apache-2.0.
Continue	commit 5522c6f44ca0 [70]	Open IDE/agent framework; Apache-2.0.
Goose	commit f1d32e8dbbdf [71]	Open extensible agent; Apache-2.0.
Roo Code	commit b867ec914575 [72]	Open editor agent; Apache-2.0.
ACP	commit 5acaa3acdall [73]	Agent-client interoperability protocol; Apache-2.0; not a repository authority system.
MCP	commit 57ac4a2ec742 [74]	Model-tool/context interoperability protocol; MIT; not a repository authority system.

APPENDIX B

EVIDENCE MANIFEST AND PUBLICATION PREFLIGHT

The machine source is `docs/paper/evidence.json`; its schema identifier is `bullet.paper-evidence.v1`. Eight required groups record the snapshot ID, UTC time and source epoch; exactly four repository commit/tree/clean subjects; tool versions; full commands with exit code, semantic outcome, and limitation; shared result values; the fixed four-term maturity vocabulary; dated external pins; and hashes of evidence-critical inputs. The build validates that object and emits `evidence.generated.tex`, which both documents consume. Gate counts, repository abbreviations, audit results, and absent transaction status are therefore not manually duplicated.

PDFs do not appear in their own input manifest: embedding a PDF hash in source that changes the PDF creates a circular fixed-point problem. Successful preflight reports final hashes as output receipts. The Hub commit has a related self-reference boundary. The snapshot binds the clean evidence subject before publication; a later Hub head is accepted only when it descends from that subject, every intervening path belongs to the enumerated paper/tooling set, and hashed evidence inputs are unchanged. Kernel, BulletGit, and Portal must match exactly. This narrow exception cannot lend old receipts to code changes.

Strict preflight rejects a dirty capture or checkout, false maturity, repository/commit/tree mismatch, unexpected Hub change, stale input hash, undefined/multiple citation or reference, duplicate or unused bibliography entry, and an external reference without URL/DOI and access date. It builds twice under one `SOURCE_DATE_EPOCH`, UTC timezone, deterministic PDF metadata, and locale; both PDF hashes must repeat. Layout/PDF checks require no material overfull box, US Letter, title/author/subject/keywords, embedded fonts, 12–14 paper pages, and exactly four brief pages. Every rendered page is then inspected because no log parser proves that a table, URL, vector arrow, or final column is visually legible.

This is publication assurance, not a product gate. A byte-identical paper can faithfully report a broken product, and a correct component can coexist with a failed paper check. The guarantee is narrower: both documents state the same machine-backed facts about the same subjects and common accidental overclaim fails closed.

REFERENCES

- [1] J. Gray, “The transaction concept: Virtues and limitations,” in *Proceedings of the 7th International Conference on Very Large Data Bases*, 1981, pp. 144–154. [Online]. Available: <https://www.vldb.org/conf/1981/P144.PDF>
- [2] J. Gray and A. Reuter, *Transaction Processing: Concepts and Techniques*. Morgan Kaufmann, 1993. [Online]. Available: <https://www.sciencedirect.com/book/9781558601901/transaction-processing>
- [3] P. Helland, “Idempotence is not a medical condition,” *Communications of the ACM*, vol. 55, no. 5, pp. 56–65, 2012. [Online]. Available: <https://doi.org/10.1145/2160718.2160734>
- [4] Bullet Farm Maintainers, “Architecture overview,” Bullet Farm source, `docs/architecture/overview.md`, 2026, last reviewed 2026-08-25; normative orientation, pre-release.
- [5] Gas Town Maintainers, “Gas town v1.2.1,” 2026, mIT; peeled commit 319d33a91b2deca59bba6dd26be6b9daf8eaacf6. [Online]. Available: <https://github.com/gastownhall/gastown/releases/tag/v1.2.1>
- [6] Gas City Maintainers, “Gas city v1.4.1,” 2026, mIT; peeled commit 58ef17e3bd685fd5cf7f21286277b208d3324590. [Online]. Available: <https://github.com/gastownhall/gascity/releases/tag/v1.4.1>
- [7] DeepSeek AI, “Deepseek harness dsh-v0.1.1-rc.2,” 2026, mIT; commit b150a551b8d465e31e418e1b2eaf5e79bbb7d28e. [Online]. Available: <https://github.com/deepseek-ai/DeepSeek-Harness/releases/tag/dsh-v0.1.1-rc.2>
- [8] Omnigent Maintainers, “Omnigent v0.10.0,” 2026, apache-2.0; commit 40755dd8dddb07e1eb6e4055d1d9936e184ceb9b. [Online]. Available: <https://github.com/omnigent-ai/omnigent/releases/tag/v0.10.0>
- [9] OpenAI, “Codex cli,” 2026, official OpenAI documentation; proprietary behavior is black-box limited. [Online]. Available: <https://developers.openai.com/codex/cli>
- [10] —, “Codex cloud,” 2026, official OpenAI documentation; hosted behavior is black-box limited. [Online]. Available: <https://developers.openai.com/codex/cloud>
- [11] Bullet Farm Maintainers, “Bullet paper evidence snapshot, `bullet.paper-evidence.v1`,” Bullet Farm source, `docs/paper/evidence.json`, 2026, shared machine-readable input for the paper and executive brief.
- [12] —, “Release evidence contract and current gate inventory,” Bullet Farm source, `docs/release.md`, 2026, status BLOCKED; last reviewed 2026-08-25.
- [13] Git Project, “git-worktree documentation,” 2026, official linked-worktree contract; Git documentation license applies. [Online]. Available: <https://git-scm.com/docs/git-worktree>
- [14] Gas Town Maintainers, “Issue 4094: Background checkout at town root deletes operational files,” 2026, open when observed; reporter states affected version 1.1.0. [Online]. Available: <https://github.com/gastownhall/gastown/issues/4094>
- [15] Gas City Maintainers, “Issue 1181: Parallel agents share the rig path,” 2026, open when observed. [Online]. Available: <https://github.com/gastownhall/gascity/issues/1181>
- [16] —, “Issue 5511: Timeout retry creates duplicate workflows,” 2026, open when observed; reporter states affected version v1.4.1. [Online]. Available: <https://github.com/gastownhall/gascity/issues/5511>
- [17] —, “Issue 1052: Concurrent sessions execute the same bead,” 2026, closed via issue 2985; reporter states affected version 0.15.2. [Online]. Available: <https://github.com/gastownhall/gascity/issues/1052>

- [18] Gas Town Maintainers, “Issue 3584: Dead polecats go undetected — no notification when a polecat session dies mid-work,” 2026, opened 2026-04-10; closed when observed; version not stated. [Online]. Available: <https://github.com/gastownhall/gastown/issues/3584>
- [19] —, “Issue 4737: mol-deacon-patrol cleanup dog-liveness check can never succeed, guarded only by an accident of worktree layout,” 2026, opened 2026-08-19; open when observed; reporter states affected version 1.1.0. [Online]. Available: <https://github.com/gastownhall/gastown/issues/4737>
- [20] —, “Issue 4091: Daemon refinery aliveness check trusts stale pid files — dead refineries never respawned,” 2026, opened 2026-05-21; open when observed; version not stated. [Online]. Available: <https://github.com/gastownhall/gastown/issues/4091>
- [21] M. Kleppmann, “How to do distributed locking,” 2016, discussion of fencing tokens checked by the storage service. [Online]. Available: <https://martin.kleppmann.com/2016/02/08/how-to-do-distributed-locking.html>
- [22] Bullet Farm Maintainers, “Bullet farm safety-complete v1 closure plan,” Bullet Farm source, `docs/assurance/v1-closure-plan.md`, 2026, authoritative documentation plan; runtime and generated evidence win.
- [23] —, “Bulletgit architecture and trust boundary,” BulletGit source, `docs/architecture.md`, 2026, component architecture; positive online authority remains blocked.
- [24] —, “Committed gate-0 policy snapshot,” Bullet Farm source, `policy/v1alpha1/policy.json`, 2026, v1alpha1, generation 1, live admission disabled, evolutionary authority disabled; hashed in the evidence manifest.
- [25] M. Salaün, “Landlock: Unprivileged access control,” Linux kernel documentation, 2025. [Online]. Available: <https://docs.kernel.org/userspace-api/landlock.html>
- [26] Linux Kernel Contributors, “Seccomp BPF (SECure COMPUting with filters),” Linux kernel documentation, 2025. [Online]. Available: https://docs.kernel.org/userspace-api/seccomp_filter.html
- [27] Linux man-pages Project, “namespaces(7) — overview of Linux namespaces,” 2025. [Online]. Available: <https://man.7.org/linux/man-pages/man7/namespaces.7.html>
- [28] containers Contributors, “bubblewrap: Low-level unprivileged sandboxing tool,” 2026, user-namespace sandbox used by Flatpak and by Omnigent’s native Linux wrappers. [Online]. Available: <https://github.com/containers/bubblewrap>
- [29] G. Hoare, “Not rocket science: The bors merge bot,” 2014. [Online]. Available: <https://graydon2.dreamwidth.org/1597.html>
- [30] GitHub, “Managing a merge queue,” 2026. [Online]. Available: <https://docs.github.com/en/repositories/configuring-branches-and-merges-in-your-repository/configuring-pull-request-merges/managing-a-merge-queue>
- [31] Bullet Farm Maintainers, “Invariant registry contract,” Bullet Farm source, `docs/assurance/invariant-registry.md` and generated registry JSON, 2026, machine-validated crosswalk; planned rows are not implemented evidence.
- [32] —, “Evolutionary multi-agent control,” Bullet Farm source, `docs/architecture/evolutionary-control.md`, 2026, normative V1 design; runtime coverage incomplete; not a claim that any evolutionary property is implemented.
- [33] J.-B. Mouret and J. Clune, “Illuminating search spaces by mapping elites,” arXiv:1504.04909, 2015. [Online]. Available: <https://arxiv.org/abs/1504.04909>
- [34] J. K. Pugh, L. B. Soros, and K. O. Stanley, “Quality diversity: A new frontier for evolutionary computation,” *Frontiers in Robotics and AI*, vol. 3, p. 40, 2016. [Online]. Available: <https://doi.org/10.3389/frobt.2016.00040>
- [35] K. Jamieson and A. Talwalkar, “Non-stochastic best arm identification and hyperparameter optimization,” in *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics (AISTATS), PMLR 51*, 2016. [Online]. Available: <https://proceedings.mlr.press/v51/jamieson16.html>
- [36] L. Li, K. Jamieson, A. Rostamizadeh, E. Gonina, M. Hardt, B. Recht, and A. Talwalkar, “A system for massively parallel hyperparameter tuning,” arXiv:1810.05934; *Proceedings of Machine Learning and Systems (MLSys)* 2020, 2020, asynchronous successive halving (ASHA). [Online]. Available: <https://arxiv.org/abs/1810.05934>
- [37] E. Alba and M. Tomassini, “Parallelism and evolutionary algorithms,” *IEEE Transactions on Evolutionary Computation*, vol. 6, no. 5, pp. 443–462, 2002, island models. [Online]. Available: <https://doi.org/10.1109/TEVC.2002.800880>
- [38] L. Li, W. Chu, J. Langford, and R. E. Schapire, “A contextual-bandit approach to personalized news article recommendation,” in *Proceedings of the 19th International Conference on World Wide Web (WWW 2010)*, 2010. [Online]. Available: <https://arxiv.org/abs/1003.0146>
- [39] L. Li, W. Chu, J. Langford, and X. Wang, “Unbiased offline evaluation of contextual-bandit-based news article recommendation algorithms,” in *Proceedings of the 4th ACM International Conference on Web Search and Data Mining (WSDM 2011)*, 2011. [Online]. Available: <https://arxiv.org/abs/1003.5956>
- [40] Bullet Farm Maintainers, “Bounded formal model contract,” Bullet Farm source, `formal/README.md` and `formal/model-lock.json`, 2026, exactly two bounded safety models; no product-wide proof.
- [41] Y. Yu, P. Manolios, and L. Lamport, “Model checking TLA+ specifications,” in *Correct Hardware Design and Verification Methods (CHARME 1999), LNCS 1703*, 1999, pp. 54–66. [Online]. Available: https://doi.org/10.1007/3-540-48153-2_6
- [42] D. J. Rosenkrantz, R. E. Stearns, and P. M. Lewis, “System level concurrency control for distributed database systems,” *ACM Transactions on Database Systems*, vol. 3, no. 2, pp. 178–198, 1978, wound-wait and wait-die. [Online]. Available: <https://doi.org/10.1145/320251.320260>
- [43] Bullet Farm Maintainers, “Bullet farm opportunity workplan,” Bullet Farm source, `docs/workplan.md`, 2026, non-authoritative backlog; V1 closure plan wins.
- [44] Gas Town Maintainers, “Gas town readme at v1.2.1,” 2026, pinned public architecture and operating contract. [Online]. Available: <https://github.com/gastownhall/gastown/blob/319d33a91b2deca59bba6dd26be6b9daf8eac6f/README.md>
- [45] Gas City Maintainers, “Gas city repository at v1.4.1,” 2026, pinned source and documentation tree. [Online]. Available: <https://github.com/gastownhall/gascity/tree/58ef17e3bd685fd5cf7f21286277b208d3324590>
- [46] DeepSeek AI, “Deepseek harness architecture at dsh-v0.1.1-rc.2,” 2026, pinned Cordis plugin and session-log documentation. [Online]. Available: <https://github.com/deepseek-ai/DeepSeek-Harness/tree/b150a551b8d465e31e418e1b2eaf5e79bbb7d28e>
- [47] —, “Experimental agent teams documentation,” 2026, pinned experimental roster, task-board, mailbox, and scope contract. [Online]. Available: <https://github.com/deepseek-ai/DeepSeek-Harness/tree/b150a551b8d465e31e418e1b2eaf5e79bbb7d28e>
- [48] Omnigent Maintainers, “Omnigent repository at v0.10.0,” 2026, pinned runner, sandbox, collaboration, and proxy documentation. [Online]. Available: <https://github.com/omnigent-ai/omnigent/tree/40755dd8dddb07e1eb6e4055d1d9936e184ceb9b>
- [49] —, “Omnigent l7 credential proxy,” 2026, pinned l7 egress and request-transformation documentation. [Online]. Available: <https://github.com/omnigent-ai/omnigent/tree/40755dd8dddb07e1eb6e4055d1d9936e184ceb9b>
- [50] —, “Omnigent contextual policy contract,” 2026, pinned ALLOW, DENY, and ASK policy vocabulary. [Online]. Available: <https://github.com/omnigent-ai/omnigent/tree/40755dd8dddb07e1eb6e4055d1d9936e184ceb9b>
- [51] —, “Omnigent polly documentation,” 2026, pinned parallel-worktree and cross-vendor review material. [Online]. Available: <https://github.com/omnigent-ai/omnigent/tree/40755dd8dddb07e1eb6e4055d1d9936e184ceb9b>
- [52] —, “Issue 1694: Polly parallel-worktree session failures,” 2026, open when observed; reporter states affected version 0.2.0. [Online]. Available: <https://github.com/omnigent-ai/omnigent/issues/1694>
- [53] J. Jena and S. Dalal, “The idempotency-key HTTP header field,” Internet-Draft draft-ietf-httpapi-idempotency-key-header-07, IETF HTTPAPI Working Group (expired), 2025. [Online]. Available: <https://datatracker.ietf.org/doc/draft-ietf-httpapi-idempotency-key-header/>
- [54] Gas Town Maintainers, “Issue 742: settings.json written to wrong directory,” 2026, closed when observed; contains no 687-respawn claim. [Online]. Available: <https://github.com/gastownhall/gastown/issues/742>
- [55] Omnigent Maintainers, “Issue 725: External changes missing in non-git workspace,” 2026, open when observed. [Online]. Available: <https://github.com/omnigent-ai/omnigent/issues/725>
- [56] —, “Pull request 4318: Revert approval attribution stack,” 2026, merged 2026-08-07 and listed as a v0.10.0 breaking change. [Online]. Available: <https://github.com/omnigent-ai/omnigent/pull/4318>
- [57] Anthropic, “Claude code overview,” 2026, official documentation; proprietary behavior is black-box limited. [Online]. Available: <https://docs.anthropic.com/en/docs/claude-code/overview>
- [58] Cursor, “Cursor agent documentation,” 2026, official documentation; proprietary behavior is black-box limited. [Online]. Available: <https://docs.cursor.com/agent/overview>

- [59] GitHub, “About github copilot coding agent,” 2026, official documentation; hosted behavior is black-box limited. [Online]. Available: <https://docs.github.com/en/copilot/concepts/agents/coding-agent/about-coding-agent>
- [60] Google, “Jules documentation,” 2026, official documentation; hosted behavior is black-box limited. [Online]. Available: <https://jules.google/docs/>
- [61] Cognition, “Devin documentation,” 2026, official documentation; hosted behavior is black-box limited. [Online]. Available: <https://docs.devin.ai/>
- [62] Amazon Web Services, “Kiro documentation,” 2026, official documentation; proprietary behavior is black-box limited. [Online]. Available: <https://kiro.dev/docs/>
- [63] Google, “Gemini cli at commit 812f7a2bcf20,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/google-gemini/gemini-cli/tree/812f7a2bcf20b6e80e2e50c3c8fa8e26567bc1e8>
- [64] OpenHands Maintainers, “Openhands at commit 150e76046db0,” 2026, mIT; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/All-Hands-AI/OpenHands/tree/150e76046db026dd944df0506642dc9b7b99391e>
- [65] —, “Agent canvas overview,” 2026, official project documentation observed with OpenHands commit 150e76046db0. [Online]. Available: <https://docs.openhands.dev/openhands/usage/agent-canvas>
- [66] SWE-agent Maintainers, “Swe-agent at commit 3ea751c087f3,” 2026, mIT; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/princeton-nlp/SWE-agent/tree/3ea751c087f32b16e039a2233dd6eecefc325d5>
- [67] mini-SWE-agent Maintainers, “mini-swe-agent at commit 25941c89cfbc,” 2026, mIT; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/SWE-agent/mini-swe-agent/tree/25941c89cfbc91eb40b3f8756348c91d9977d57e>
- [68] Aider Maintainers, “Aider at commit 5dc9490bb35f,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/Aider-AI/aider/tree/5dc9490bb35f9729ef2c95d00a19ccd30c26339c>
- [69] Cline Maintainers, “Cline at commit 4f5f2384076e,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/cline/cline/tree/4f5f2384076e1eef461f53c8eb7d0ece6e0bb552>
- [70] Continue Maintainers, “Continue at commit 5522c6f44ca0,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/continuedev/continue/tree/5522c6f44ca0ac3528b37244818bfa39b5af470>
- [71] Block, “Goose at commit f1d32e8dbbdf,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/block/goose/tree/f1d32e8dbbdf4c7265178904f00283af110ada1c>
- [72] Roo Code Maintainers, “Roo code at commit b867ec914575,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/RooCodeInc/Roo-Code/tree/b867ec9145750d0ae1ff7f02d35406e9bf2a0b16>
- [73] Agent Client Protocol Maintainers, “Agent client protocol at commit 5acaa3acda11,” 2026, apache-2.0; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/agentclientprotocol/agent-client-protocol/tree/5acaa3acda11e29d523193bfb3401c8497f5a90e>
- [74] Model Context Protocol Maintainers, “Model context protocol at commit 57ac4a2ec742,” 2026, mIT; exact commit observed 2026-08-25. [Online]. Available: <https://github.com/modelcontextprotocol/modelcontextprotocol/tree/57ac4a2ec742e0cb7622d899b0f5d3bcf769fd69>

APPENDIX C

RED-TEAM SELF-REVIEW

Ten attacks a hostile reviewer will run, each answered or conceded.

- 1) “Your issues do not say what you say.” Conceded for an earlier draft that attributed a respawn storm and a

“687 respawns” figure to issue 742; that issue concerns a settings path and the figure has no source. The liveness class now cites issues 3584, 4737, and 4091 with state, date, and pinned corroboration (Table VII). Residual: trackers move.

- 2) “*Component receipts are not a product.*” Conceded. Nothing here exceeds Component-proved; 26 of 26 release gates are blocked and the connected transaction is absent.
- 3) “*Two TLA+ models prove the system.*” Conceded. Bounded safety only, deadlock checking off, one invariant true by assignment, one defined but unchecked, no wound-wait, Git, forge, or provider state.
- 4) “*Fencing tokens, idempotency keys, sandboxes, and merge queues with new names.*” Largely conceded and cited. The claim is compositional: the same subject and principal at every boundary. Its worth is the deferred Stage-2 question.
- 5) “*Providers propose is unenforceable.*” Partly conceded. Only a validated proposal crosses the writer boundary, but the provider filesystem and process sandbox is Designed; only the network-namespace egress component exists, and no transaction has composed admission, proposal, and apply.
- 6) “*An evolutionary paper with no evolutionary code.*” Conceded. Section III-F is design; the policy pins `evolutionary_authority=false`; the program is post-V1 and gated on the connected transaction.
- 7) “*A single-user loopback threat model is trivial.*” Narrow by design: the adversaries are the model, a concurrent incarnation, a retrying supervisor, hostile repository content, and a promoting projection, which are the corpus failures. Multi-tenant and hostile-host threats are named out of scope, not implied.
- 8) “*An agent can flip live admission.*” Rejected. The schema rule is enforced by the Hub validator and mirrored by the Kernel loader; live admission needs generation ≥ 2 and a registered runner key; real-mode nightly refuses without an operator-ratified policy path; an environment token was rejected as admission.
- 9) “*A security paper with 29 hard findings.*” Conceded. The figures are generated from the manifest, dated, and block release by the family’s own rule; no security quality is claimed.
- 10) “*The comparison is a cartoon of featureful systems.*” Rejected. Operational capability and authority/integrity are separate matrices; Not documented is never converted into “No”; each deep study credits what the system contributes; no superiority language precedes the matched benchmark.